

Part 1 Comments:

- The variables did not need data type conversion. And there were no duplicate rows.
- The max values for MedInc, AveRooms, AveBedrms, Population, and AveOccup are oddly high. Therefore, I ran the same models on the original dataset and then again on the dataset but with the outliers (anything outside of Q1 and Q3) removed. I found that despite the model's metrics performed worse on the surface (the numbers were lower), there were fewer false positives and false negatives on the dataset with the outliers removed.
- Visualizing the dataset, I found that the dataset is very skewed to the right. In addition, I noticed that the houses with prices above the median are symmetrical; this could possibly explain why the model metrics are almost the same for each model (i.e., precision, recall, and f1-score are close in number for each model). However, after removing the outliers, the dataset is generally a normal distribution.

Which techniques did you use to train the models?

I used four supervised learning models: K-Nearest Neighbors, Decision Tree, Random Forest, and AdaBoost. Before I trained the model, I cleaned the dataset by checking and removing the duplicates and graphing the data to ensure it was correct. In order to train the model, I first split the dataset into training and testing sets using a stratified split to maintain the proportion of the target variable. In order to improve the model, I then standardized features using a StandardScaler. I also removed outliers to try to improve the model even more. Finally, I fit each model to the training data using their respective fit() methods.

Explain any techniques used to optimize model performance?

To optimize performance, I standardized features using StandardScaler so that distance-based algorithms like KNN perform optimally. In addition, I tested different test/train splits to find the best split to ensure that the class distribution of the target variable is preserved in both the training and testing sets, reducing sampling bias. Finally, I removed outliers to try to optimize the model performance even further.

Compare the performance of all models to predict the dependent variable?

KNN: This model generally performed well when features were standardized, but its performance can be sensitive to the choice of neighbors and local data variations. In addition, I noticed that because I tested so many neighbors (1-21), this model took a noticeably larger amount of time to run.

Decision Tree: This model provided an easily interpretable model. However, it can be prone to overfitting if not properly pruned or tuned. I noticed that despite the high performance, this model had the most false positives and false negatives.

Random Forest: This model performed the best overall. Random forest had the best performance by averaging multiple trees, which reduced overfitting and was able to deliver better performance.

AdaBoost: This model also performed pretty well by sequentially correcting misclassifications, though it can be sensitive to noisy data. I think this performed well on this dataset because this dataset was pretty clean.

Which model would you recommend to be used for this dataset?

Based on the results, I would recommend using the Random Forest classifier for this dataset. Random Forest demonstrated the most balanced performance, with high accuracy and strong generalization capabilities. It effectively handles non-linear relationships and avoids overfitting better than a single Decision Tree. Additionally, this model had the lowest number of false positives and false negatives.

For this dataset, which metric is more important, why?

For this dataset, I think recall is the most important metric because it ensures that we correctly identify as many high-value houses as possible. In real estate pricing, missing a high-value house (false negatives) could result in undervaluing properties, which might lead to financial loss. However, if both false positives and false negatives need to be balanced, the F1-score would also be a suitable metric since it considers both precision and recall.